

Optimizing warnings for hydrologic ensemble prediction systems for improved decision making

Jesús Casado-Rodríguez,^a Corentin Carton de Wiart,^b Stefania Grimaldi,^a Ervin Zsoter,^b Calum Baugh,^b Nina Bosshard,^c Michaela Mikulickova,^d Ilias Pechlivanidis,^c Christel Prudhomme,^b and Peter Salamon^a

^a *European Commission - Joint Research Centre, Ispra (Italy)*

^b *European Centre for Medium-Range Weather Forecasts, Reading (UK)*

^c *Swedish Meteorological and Hydrological Institute, Norrköping (Sweden)*

^d *Slovak Hydrometeorological Institute, Bratislava (Slovakia)*

Corresponding author: Jesús Casado-Rodríguez, jesus.casado-rodriguez@ec.europa.eu

11 ABSTRACT: Flood Early Warning Systems (FEWS) rely on hydrological simulations driven by
12 Numerical Weather Prediction (NWP) models, both of which are inherently uncertain. Ensemble
13 Prediction Systems (EPS) address these uncertainties by generating multiple future scenarios.
14 Discharge forecasts from EPS are converted into flood warnings by applying a set of criteria
15 whose definition is crucial for the skill of the system. While meteorological and hydrological
16 models are under continuous development, these warning criteria lack continuous evaluation. In
17 this paper, we use discharge simulations from the European Flood Awareness System (EFAS)
18 to assess the skill of four NWPs —probabilistic and deterministic—, and search for methods to
19 combine them into a grand ensemble to enhance overall skill. Additionally, we evaluate the effects
20 of the current warning criteria —probability threshold and persistence— and determine their
21 optimal values. Our results confirm that probabilistic NWPs outperform deterministic models in
22 terms of flood warning skill. We demonstrate that the persistence criterion should be omitted from
23 EPS as it hinders skill. Using a single probabilistic NWP with optimal warning criteria enhances
24 the current EFAS skill by 6.6% (measured in terms of f-score), and building a grand ensemble can
25 further increase skill by 4.2%. We identify two effective grand ensemble methods —member-based
26 and skill-based— and discuss their advantages and drawbacks. In conclusion, our study presents a
27 methodology for evaluating and refining the skill of multimodel FEWS, stressing the critical role
28 of tailoring the flood warning criteria to the end-users.

29 SIGNIFICANCE STATEMENT: This study analyses methodologies to improve the skill of the
30 notifications issued by a state-of-the-art flood warning system. We have analysed the skill of four
31 meteorological models and the effects of several warning criteria. We have identified an optimal
32 blend of meteorological models and warning criteria that boosts the system skill up to 10%. The
33 methods developed in this research are adaptable to other flood warning systems, and the refined
34 criteria will be implemented in the operational system.

35 **1. Introduction**

36 Floods are the natural disaster that cause the biggest economical losses in Europe (WMO 2021).
37 In the European Union, flood damages amounted to a total cost of approximately EUR 280 billion in
38 the period 1980-2022, representing 43% of all meteorological hazards (EEA 2023). The historical
39 data show an increasing tendency in flood damages over Europe, and climate projections indicate
40 an exacerbation of this trend with increasing frequency and severity of extreme events (EEA 2023;
41 Paprotny et al. 2018; Alfieri et al. 2017; Caretta et al. 2023). In this scenario, Flood Early Warning
42 Systems (FEWS) emerge as one of the multiple adaptation measures to be taken. The magnitude
43 of the flood that hit Germany and Belgium in July 2021, the costliest in the last 40 years, was a
44 reminder of the importance of FEWS and their need for continuous improvement of forecasts and
45 decision making (EEA 2023).

46 FEWS rely on meteorological forecasts generated by Numerical Weather Prediction (NWP)
47 models (Pagano et al. 2014). Over the past thirty years, Ensemble Prediction Systems (EPS) have
48 been developed within the realm of NWP models to address uncertainties in the initial conditions.
49 EPS provide equiprobable forecasts at any given time and location, enabling the quantification
50 and dissemination of probabilistic forecasts (Cloke and Pappenberger 2009). In comparison to
51 deterministic models, EPS have demonstrated enhanced consistency across consecutive forecasts
52 and improved precipitation forecasts at longer lead times (Buizza 2008; Molteni et al. 1996).

53 The success of EPS in meteorology led to the adoption of this paradigm in flood forecasting (De
54 Roo et al. 2003; Bartholmes and Todini 2005; Pappenberger et al. 2005; Roulin and Vannitsem
55 2005; Diomedea et al. 2008; Yang et al. 2020). Hydrological Ensemble Predictions Systems (HEPS)
56 propagate the uncertainties in the meteorological forecasts by forcing hydrological models with
57 the members of a meteorological ensemble. Although this approach is the most common, it only

58 considers the uncertainty in the initial meteorological condition, but does not include other sources
59 of uncertainty, such as model structure or model parameterization (Cloke and Pappenberger 2009).

60 Issuing weather-related warnings, such as floods, is a delicate decision as it might result in
61 significant costs, either due to unpreparedness if an event is missed or the mobilization of resources
62 in false alarms. Weather forecasting systems often issue a high number of false alarms, despite the
63 potential risk of eroding trust in the system, known as the "cry wolf" effect (Bouttier and Marchal
64 2024; LeClerc and Joslyn 2015). This latter research showed that an increase in false alarms
65 diminishes the credibility of the system, leading end-users to disregard precautionary measures.
66 However, they also demonstrated that incorporating uncertainty estimates in warning messages can
67 assist users in making informed decisions. Therefore, HEPS have the potential to enhance both
68 the skill and trustworthiness of flood warning system, if particular attention is given to the balance
69 between false alarms and missed events.

70 The major flood events in the Oder River in 1997 and the Elbe River in 2002 triggered the
71 development of the European Flood Awareness system (EFAS), a continental scale FEWS that
72 has been operational since 2012 as part of the Copernicus Emergency Management Service of
73 the European Commission (Bartholmes et al. 2009; Thielen et al. 2009a). EFAS is an operational
74 system that monitors and forecasts floods in an extensive European domain. Its primary objective is
75 to enable proactive measures in anticipation of significant flood events. EFAS is a complementary
76 system to national or regional counterparts, with a specific focus on trans-boundary rivers and
77 medium-range forecasts.

78 With the release of version 4.0 in October 2020, EFAS has been generating forecasts twice a
79 day, with a 6-hour resolution and lead times extending up to 10 days. It couples four NWP of
80 diverse nature —deterministic and probabilistic— with the distributed, process-based hydrolog-
81 ical model LISFLOOD-OS (De Roo et al. 2000; Burek et al. 2013). The simulation produces a set
82 of water fluxes and state variables, from which river discharge is utilized to issue flood warnings.
83 EFAS employs a threshold exceedance approach for issuing warnings, where continuous discharge
84 time series are converted into binary series of exceedance or not-exceedance over a discharge
85 threshold (Ramos et al. 2007; Thielen et al. 2009a; Bartholmes et al. 2009). These works devel-
86 oped a warning system based on two criteria: the *probability threshold* that defines when there

87 is significant risk of flooding in a particular forecasts, and the *persistence* of that signal across
88 consecutive forecasts.

89 The criteria applied in EFAS to issue flood warnings have changed over the past years driven
90 by experience of the forecasters and feedback from users. Previous studies have examined EFAS
91 hydrological performance (Alfieri et al. 2014), its warning skill (Bartholmes et al. 2009; Thielen
92 et al. 2009b), and the response of end users (Demeritt et al. 2013). However, the continuous
93 advances in NWP performance (Bauer et al. 2015), as well as in the resolution and representation
94 of hydrological processes in EFAS (JRC 2020, 2023), requires a new evaluation of the warning
95 criteria to harness the potential of the system.

96 The objective of this analysis is to assess the skill of a continental flood forecasting system
97 like EFAS in its current state, and explore new warning criteria that can maximize that skill. In
98 particular, the first goal is to determine the value of combining several NWPs of different nature in a
99 single system, specifically seeking the appropriate combination method to build a grand ensemble.
100 The second goal is to optimise the two warning criteria, analysing the interplay between them and
101 reconsidering whether the persistence criterion is meaningful in a purely probabilistic approach.
102 The third and final goal is to analyse the relationship between warning skill and catchment area to
103 discern if the newer, higher resolution models are skilful in smaller catchments.

104 To address our research questions, we conducted an analysis of EFASv4 discharge simulations
105 for the period 2020-2023. Section 2 enumerates the dataset used in the analysis. Section 3 explains
106 the two experiments that analyse the skill of individual NWPs and different combinations of those
107 models, the criteria optimization process, and the skill metrics we targeted. Section 4 presents
108 first the skill of individual NWP, followed by that of the combination methods, and ends with an
109 exploration of the relationship between warning skill and catchment area. Section 5 discusses
110 the importance of the spatio-temporal framework and target metric in skill assessments, the trade-
111 offs between the warning criteria, the benefits of combining NWP models, and the challenges in
112 selecting a combination method.

113 **2. Data**

114 The data used in this analysis is taken from the operational setup of the EFASv4 (JRC 2020).
115 EFAS generates a hydrological forecast twice a day (00 and 12 UTC) forced by the four NWPs

TABLE 1. Characteristics of the NWP models used in EFASv4.

Model	Provider	Acronym	Max. lead time	No. ensembles	Spatial resolution
COSMO-LEPS		COS	5.5 days	20	~ 7 km
ICON-EU/ICON	DWD	DWD	7 days	1	~ 6.5-13 km
HRES	ECMWF	HRES	10 days	1	~ 9 km
ENS	ECMWF	ENS	10 days	51	~ 18 km

outlined in Table 1. EFAS integrates forecasts from probabilistic models (COS and ENS) and deterministic models (DWD and HRES) to issue flood warnings.

The study period spans from the release of EFASv4 in October 2020 until June 2023. We use EFAS forecasts as the discharge prediction and the EFAS reanalysis as a proxy of discharge observations. Using reanalysis as ground truth addresses challenges posed by limited hydrological observations (e.g. gaps in the observed time series, difference in recorded period, etc), and removes model representation and calibration errors from the analysis (Pappenberger et al. 2015; Cloke et al. 2017).

In contrast to the original EFAS skill assessment (Bartholmes et al. 2009), the current analysis focuses solely on real gauging stations rather than all river cells in the model, which would have caused an autocorrelation issue. From the approximately 4000 gauging stations in the EFAS database, we filtered stations based on two conditions: the contributing area must exceed 500 km², and the hydrological performance, in terms of the modified KGE (Kling-Gupta efficiency) (Kling et al. 2012; Gupta et al. 2009; Knoben et al. 2019), must not fall below 0.50. The lower bound on catchment was set due to the relatively coarse spatial resolution of NWP and EFAS temporal resolution (6 h), which limits hydrological performance in small catchments. The condition based on hydrological performance ensures that the hypothesis of using model simulations as proxies for observations is valid. Ultimately, 1979 gauging stations were included in the analysis, accounting for a total of 1683 events. As explained in section d, the optimization of the warning criteria was conducted on a subset of 1239 stations with a contributing area larger than 2000 km² (874 events), while the entire set was used to analyse the influence of catchment area on warning skill. Appendix A illustrates the distribution of the selected stations and the number of observed events.




FIG. 1. Overview of the skill assessment framework. a) Hydrographs comparing observed and forecasted discharge with the discharge associated to the 5-year return period (Q_5). b) Matrix of total exceedance probability—white dots identify cells that meet the warning criteria, whereas black rectangles delineate the location of the forecast hydrograph in panel (a)—, timeline of predicted events for that specific criteria, and timeline of observed events where hits are highlighted in green and misses in orange. c) Confusion matrix based on the previous warning criteria. d) Evaluation of the f-score across all combinations of probability threshold and persistence—the black square highlights the skill score corresponding to the warning criteria in panel (b)—.

3. Methods

The skill assessment comprises four steps, as depicted schematically in Figure 1: (a) transforming both observed and forecasted discharge time series into corresponding time series of probability of exceeding a specific return period; (b) building an exceedance probability matrix that consolidates the overlapping forecasts, followed by identifying predicted events through the application of the warning criteria; (c) creating the contingency table, (d) and evaluating the skill of the system across all possible combinations of the warning criteria.

152 *a. Predicting flood events*

153 The warning process starts by converting the continuous time series of the target variable into
154 time series of exceedance or non-exceedance over a predefined threshold. Since forecasts are issued
155 more frequently than the forecast horizon, they overlap. To address this, the ensemble forecasts
156 are restructured into an exceedance probability matrix that connects each time step to increasing
157 lead times, as illustrated in panel (b) of Figure 1. Within this matrix, an individual forecast is
158 represented by a diagonal (black rectangles). The uppermost rows represent shorter lead times,
159 where the probability of accurate prediction generally increases. For instance, in the example in
160 Figure 1, the flood event occurring in July 16 was predicted with 50% probability five days in
161 advance, increasing close to certainty within two days of the event.

162 The derivation of the binary time series of observed events is straight forward by application
163 of the threshold over the discharge reanalysis. Conversely, deriving the binary time series of
164 predicted events involves applying the warning criteria to the matrix of probability of exceedance.
165 The *probability threshold* sets the minimum value of the exceedance probability at which a high
166 risk of flooding is considered. Additionally, the *persistence* criterion removes false alarms caused
167 by the erratic behaviour of some NWP (Bartholmes et al. 2009). This criterion aims to imitate the
168 behaviour of a forecaster, who would wait to send the warning until consecutive forecasts predict
169 the event. Panel (b) in Figure 1 illustrates the application of these two warning criteria on the
170 matrix of exceedance probability. For every time step (columns in the total probability matrix), the
171 white dots identify the lead times that meet the criteria. A flood event is predicted at a particular
172 time step if any of the lead times meet the criteria.

173 EFAS CURRENT WARNING CRITERIA

174 Flood warnings in EFAS are based on discharge, and the critical threshold is the discharge
175 associated with the 5-year return period (Q_5), computed by fitting a Gumbel distribution to the
176 annual maxima in the EFASv4 long-term discharge simulation (1990-2022). While alternative
177 alert systems might rely on disparate variables, such as precipitation (Bouttier and Marchal 2024)
178 or river stage (Nevo et al. 2022), and employ different thresholds for defining events, the underlying
179 methodology remains similar.

TABLE 2. Contingency table in a binary classification. E_{obs} is the sum of observed events and E_{pred} the sum of predicted events.

		Observed		E_{pred}
		True	False	
Forecasted	True	hit	false alarm	
	False	miss	true negative	
		E_{obs}		

To issue a flood warning, five criteria must be met. EFAS warnings are exclusively issued to locations with a minimum catchment area of 2000 km² and for lead times longer than 48 hours. The first limit is based on the inferior hydrological performance in smaller catchments (Alfieri et al. 2014), while the second limit is a consequence of EFAS addressing early awareness and serving as a supplementary system to national or regional services. A warning is issued if at least one deterministic and one probabilistic NWP predict a probability of exceeding Q_5 larger than 30% over 3 consecutive forecasts (hereafter referred to as 3/3 persistence). We denote this approach of combining the NWP models as *1 deterministic + 1 probabilistic* (1D+1P). As each model is evaluated independently, the current procedure does not calculate a total probability matrix.

b. Assessing warning skill

1) CONTINGENCY TABLE

Assessing the skill of a warning system is a binary classification task that compares two time series: the predicted and observed events. The contingency table (2) is commonly used to summarize the the performance of such tasks. Given that floods are rare events, the classification task is highly imbalanced, resulting in the number of true negatives being orders of magnitude larger than any of the other terms in the contingency table. Consequently, in order to avoid overestimating skill, we must omit the true negatives in the selection of the target skill metric (Section 2).

The terms of the contingency table result from comparing the time series of predicted and observed events. We consider a hit any observed event in which at least one time step was correctly predicted. The number of misses and false alarms are calculated as the difference between the observed (E_{obs}) or predicted (E_{pred}) events and the hits, respectively.

$$\text{misses} = E_{obs} - \text{hits} \quad (1)$$

$$\text{false alarms} = E_{pred} - \text{hits} \quad (2)$$

The process is repeated for all possible combinations of the two criteria and across all stations. We tested probability threshold values ranging from 5% to 95% and six persistence values: 1/1 (no persistence), 2/4, 2/3, 2/2, 3/4 and 3/3.

2) SKILL METRICS

The highly imbalanced nature of our classification tasks is fundamental in the selection of the skill metric. Metrics such as the odds ratio and the Hanssen-Kuipers score (HK) (Hanssen and Kuipers 1965), which account for true negatives, are unsuitable for imbalanced classification and were thus excluded. Instead, we selected three metrics specific for imbalanced classification.

Recall is the proportion of observed events that are correctly predicted (eq. 3). Bartholmes et al. (2009) proved that recall is equal to HK for highly imbalanced classifications. Precision is the proportion of the predicted events that are correct (eq. 4). There is a well-known trade-off between recall and precision. Strict criteria result in fewer, more certain warnings, leading to high precision but low recall due to many missed events. Conversely, relaxed criteria increase the number of warnings, improving recall at the cost of precision due to a rise in false alarms. As a compromise, the f_β score is a metric that balances recall and precision (eq. 5). In its most common version ($\beta = 1$) it is the harmonic mean of recall and precision, granting equal importance to both terms. However, higher importance can be given to precision ($\beta < 1$), therefore limiting the amount of false alarms, or to recall ($\beta > 1$), reducing the number of misses. In our study, after suggestions from forecasters analysing EFAS on a daily basis, we selected $f_{0.8}$ as the target skill metric. This metric balances precision and recall giving a slightly higher value to the former. The idea behind is to limit the amount of false alarms, which would jeopardize the trust of the recipients of the warnings. Both recall, precision and the f_β score range from 0 to 1, being 1 their optimal value.

As a final validation metric we used bias, which represents the proportion between the number of predicted and observed events (eq. 6). Its optimal value is 1, which means that the number

of warnings is equal to the number of observed events, and it ranges from 0 to infinity. All four metrics presented here can be plotted in the Roebber diagram (Roebber 2009) (Appendix C).

$$\text{recall} = \frac{\text{hits}}{E_{obs}} = \frac{\text{hits}}{\text{hits} + \text{misses}} \quad (3)$$

$$\text{precision} = \frac{\text{hits}}{E_{pred}} = \frac{\text{hits}}{\text{hits} + \text{false alarms}} \quad (4)$$

$$f_{\beta} = \left(1 + \beta^2\right) \frac{\text{precision} \cdot \text{recall}}{\beta^2 \cdot \text{precision} + \text{recall}} \quad (5)$$

$$\text{bias} = \frac{E_{pred}}{E_{obs}} = \frac{\text{hits} + \text{false alarms}}{\text{hits} + \text{misses}} = \frac{\text{recall}}{\text{precision}} \quad (6)$$

c. Combination of NWP

One of the strengths of EFAS is the concurrent use of four NWPs, both deterministic and probabilistic. Using multiple NWPs serves two purposes: to allow national services to compare their forecasts with the corresponding EFAS simulations, and to leverage the strengths of different NWPs. In relation with this second purpose, a primary aim of this study is to develop an effective method for integrating deterministic and probabilistic forecasts. In practical terms, EFAS forecasts produce four distinct probability matrices, one for each NWP, in contrast to the single matrix depicted in panel (b) of Figure 1. The challenge lies in devising a methodology to combine these four matrices into a total exceedance probability matrix, and enhance the system's predictive accuracy (Wetterhall et al. 2013). We explored three different approaches for calculating the total exceedance probability and compared them with the approach currently used in EFAS (Section a).

The most straightforward method to compute the total probability is the *model mean* (MM); it is a simple mean over the NWP-specific exceedance probability matrices, where all models are attributed equal weight. In this approach, the single member of a deterministic model is considered significantly more important than any individual member of a probabilistic model. An alternative approach would be to assign the same weight to each member. In this approach, referred to as *member weighted* (MW), each NWP receives a weight relative to the number of members it contains, so probabilistic models prevail over deterministic ones. Neither of these approaches consider the performance of the NWP. To address this limitation, the *Brier weighted* approach (BW) assigns a weight to each model based on its probabilistic skill. The Brier score—a probabilistic

error metric that ranges from 0 to infinite, with 0 being the optimal value—is used as a measure of probabilistic skill (Brier 1950).

$$BS = \frac{1}{T} \sum_{t=1}^T (P_{obs,t} - P_{pred,t})^2 \quad (7)$$

where BS is the Brier score of a single station, T is the number of time steps, $P_{obs,t}$ is the observed probability of exceedance, and $P_{pred,t}$ is the predicted probability of exceedance at a specific time step. We discovered that the results of the BW approach are highly sensitive to how Brier scores are converted into weights. The conversion function must invert BS —as better models, with smaller BS , should receive higher weight—and should be exponential to amplify the differences among NWP. Inverse distance weighing fulfils both conditions (eq. 8). Values of the exponent p from 1 to 9 were tested, and 7 was found to be the optimal exponent.

$$w_{nwp,lt} = \frac{BS_{nwp,lt}^{-p}}{\sum_{i=1}^4 BS_{i,lt}^{-p}} \quad (8)$$

d. Optimising warning criteria

Issuing a warning involves 5 criteria: minimum catchment area, minimum lead time, combination of NWP, probability threshold and persistence. We analysed each of these criteria individually to understand its effects, and then searched for the set of criteria that optimizes the system skill.

We have conducted two main experiments to explore the benefits of combining NWP. The first experiment analyses each NWP individually with two objectives: identifying the strengths and weaknesses of each NWP, and understanding how the warning criteria affect differently probabilistic and deterministic NWPs. The second experiment compares the four combinations of NWP explained in Section c. The aim is to identify the most appropriate combination method and compare it against the most skilful NWP to assess whether the grand ensemble adds value to the system. The setup of these two experiments is similar and is explained in the following paragraphs.

For the majority of the analyses we used only stations that met the current minimum catchment area of 2000 km². With this subset of stations, we explored the evolution of skill with lead time, probability and persistence. Lead times were grouped in two ways. First, we grouped them daily to analyse the evolution of skill with lead time. Second, for the optimization of the warning criteria,

we defined 4 lead time ranges: a) shorter than 2 days, a range at which EFAS warnings cannot be issued, b) 2-5.5 days, when all NWP are available, c) 5.5-7 days, when 3 NWP are available, d) 7-10 days, when only the ECMWF forecasts are available.

The optimization of the probability threshold and persistence criteria sought the combination of these two criteria that produces the maximum $f_{0.8}$ at each of the 4 lead time ranges. To ensure a robust selection of criteria, we applied a 10-fold cross validation: 20% of the stations were left aside as a test set, while the remaining 80% were subdivided in 10 folds. We computed the skill for every combination of 9 folds and then averaged over folds. The optimal warning criteria were derived from this average-over-fold skill matrix, as shown in panel (d) of Figure 1. Subsequently, the selected criteria were evaluated on the test set to prevent overfitting.

Lastly, we explored the impact of catchment area on warning skill using the entire set of stations. We calculated the warning skill for increasing values of the minimum catchment area, ranging from 500 km² to 300,000 km², using the previously optimized criteria.

4. Results

a. Analysis of individual NWPs

1) WARNING SKILL OF INDIVIDUAL NWPs

Figure 2 exhibits the evolution of the warning skill with lead time and probability threshold for a scenario with no persistence. As a benchmark, the skill of the current warning criteria is represented with a black, solid line. Each colour line represents a different probability threshold. For the two deterministic models (DWD and HRES) all lines overlap, since the probability threshold does not affect a deterministic forecast.

The probabilistic models show overall higher skill than the deterministic. Under certain conditions, they even outperform the current warning criteria, indicating potential for enhancing the skill of EFAS. Deterministic models demonstrate skill on par with the current criteria solely for the first day forecast. As expected, warning skill degrades with lead time. This trend is consistent across all cases, although the rate of degradation depends on the probability threshold.

As we are seeking to identify a single optimal probability threshold for the warning criteria, the selection of the optimal criteria will prioritize a lead time range from 2 days to 5.5 days. This




FIG. 2. Evolution of the warning skill with daily lead times and probability threshold for each NWP model in a scenario with no persistence. Each plot represents a different NWP. As a benchmark, the black, solid line represents the skill of the current operational criteria.




FIG. 3. Evolution of the warning skill with probability threshold and persistence for each NWP model at a fixed lead time range (from 2 to 5.5 days). Each plot represents a different NWP. As a benchmark, the black cross represents the current criteria.

approach allows us to mitigate the substantial uncertainty associated with longer lead times, and focus on the lead times at which the majority of flood warnings are issued.

2) IMPACTS OF PERSISTENCE ON THE WARNING SKILL OF INDIVIDUAL NWPs

Figure 3 illustrates the evolution of skill with probability threshold and persistence for a fixed lead time range (2-5.5 days). The benchmark (current warning criteria) is now represented by a single point. As the probability threshold does not affect deterministic forecasts, the lines for DWD and HRES are horizontal.

Three out of four NWPs are able to replicate or improve the skill of the current warning criteria at certain combinations of persistence and probability. Only DWD is less skilful than the benchmark. The skill of the probabilistic models shows a trade-off between probability threshold

and persistence. Both criteria play the same role of removing false positives by taking stricter values. Consequently, stricter persistence necessitates lower probability thresholds to optimise skill, and vice versa. Nevertheless, the highest skill is consistently obtained when persistence is removed (1/1). In this scenario, there is a wide range of probability thresholds (from 40 to 90%) that perform similarly well.

Persistence does have a significant impact in the skill of deterministic models, improving dramatically with the addition of any persistence. For the most stringent persistence (3/3), the skill of the deterministic HRES matches the benchmark.

3) OPTIMAL WARNING CRITERIA FOR INDIVIDUAL NWP

The outcome of optimising the warning criteria for individual NWPs is summarised in Table 3. The results are organised by lead time ranges; for each range, the table presents the optimized criteria and the skill metrics for the available NWPs. As a benchmark, the first row in each lead time range shows the skill of the current operational criteria. The Roebber diagram in Appendix C summarizes visually these results.

The current EFAS setup (1D+1P) is negatively bias, i.e., prioritizes precision over recall. Warnings are issued with a high level of certainty, leading to a limited number of false alarms, but many events are missed. The current criteria are not optimal as it is consistently outperformed in terms of $f_{0.8}$ at any lead time by individual NWPs, suggesting that the current combination of all NWPs hinders the skill of that single best NWP.

The analysis of the results is similar for any of the four lead time ranges. The probabilistic NWPs (COS and ENS) constantly outperform in terms of $f_{0.8}$ both the benchmark and the deterministic counterparts. Even though both probabilistic NWPs show similar skill at the shortest lead time, ENS stands out as the best model at medium-range forecasts. The skill of HRES, despite being deterministic, is comparable to the benchmark, even outperforms it at the longest lead time. Overall, the optimization of individual NWPs was able to increase the warning skill by up to 6.6%.

The optimization removed the persistence criterion for the probabilistic NWPs. Its role —removing uncertain events— is now played solely by the probability threshold. The optimal value of this probability threshold is very high at short lead times and decreases with increasing lead time, being larger than the current value up to lead time 7 days. On the contrary, deterministic models

TABLE 3. Summary of the optimization of the warning criteria for individual NWP and four lead time ranges. The initial row serves as a benchmark, indicating the skill of the current warning criteria. In each lead time range, bold fonts highlight the highest skill for a specific metric.

lead time	model	probability	persistence	recall	precision	bias	$f_{0.8}$
<2 d	1D+1P	0.300	3/3	0.522	0.735	0.710	0.634
	COS	0.875	1/1	0.673	0.718	0.937	0.700
	DWD	-	2/2	0.585	0.583	1.003	0.583
	ENS	0.800	1/1	0.681	0.711	0.958	0.699
	HRES	-	2/2	0.580	0.671	0.864	0.632
2-5.5 d	1D+1P	0.300	3/3	0.413	0.618	0.668	0.518
	COS	0.725	1/1	0.387	0.687	0.560	0.518
	DWD	-	3/3	0.420	0.522	0.805	0.477
	ENS	0.650	1/1	0.416	0.725	0.574	0.562
	HRES	-	3/3	0.416	0.612	0.680	0.517
5.5-7 d	1D+1P	0.300	3/3	0.215	0.612	0.351	0.356
	DWD	-	2/2	0.291	0.358	0.813	0.328
	ENS	0.450	1/1	0.264	0.605	0.436	0.402
	HRES	-	2/2	0.273	0.429	0.636	0.351
7-10 d	1D+1P	0.300	3/3	0.120	0.625	0.192	0.237
	ENS	0.175	2/2	0.287	0.396	0.725	0.345
	HRES	-	2/2	0.244	0.336	0.726	0.293

require persistence to enhance their skill, although there is not a clear pattern between the optimal persistence and lead time.

b. Analysis of combined NWPs

1) WEIGHING MATRICES

Three of the four combination methods compute total probability, which requires assigning weights to each NWP and lead time. Figure 4 illustrates the distribution of weights among NWP models for these three combination methods.

The *model mean* (MM) attributes equal weights to each model. The changes in the weight distribution occur when the maximum lead time of a model is exceeded. This approach assigns equal importance to deterministic and probabilistic models.




FIG. 4. Weights applied to compute total probability. Every plot represents a different method, with colours representing NWPs (blue for probabilistic and orange for deterministic models).

Instead, in the *member weighted* (MW) combination, every model run is assigned an equal weight. ENS, which has a significant larger number of members, prevails over any other NWP. Even for the first 5 days, when all the models are available, ENS receives 70% of the weight. In general, this approach heavily relies on the skill of the largest ensemble, which is not necessarily the most skilful.

In the *Brier weighted* (BW) combination, we used the Brier score (6) to derive weighing factors that prioritize the more skilled models. At very short lead times, the deterministic models demonstrate skill (Figure B1), accounting for approximately a third of the total weight. As the lead time increases, the probabilistic models, particularly ENS, assume the majority of the weight. At lead times from 2 to 5.5 days, when most of the warnings are sent, the superior skill of probabilistic models is represented by a combined weight of approximately 80%.

2) WARNING SKILL OF COMBINED NWPs

Figure 5 replicates the analysis of the warning skill previously done for the individual NWP models, but for the combination methods. It exhibits the evolution of the warning skill with lead time and probability threshold for a scenario without persistence. The benchmark is ENS, the NWP that proved highest skill.

The results indicate that ENS is responsible for the majority of the skill in the grand ensembles, since only marginal gains are possible with specific combinations at certain lead times. Only two combination methods (MW and BW) are able to replicate the skill of the benchmark across all lead times. Although MM at lead time 0 achieves the highest marginal gain, this combination method




FIG. 5. Evolution of the warning skill at daily lead times and probability threshold for each of the combination methods in a scenario with no persistence. Each plot represents a different combination of NWP. As a benchmark, the black, solid line represents the skill of ENS with optimised criteria.




FIG. 6. Evolution of the warning skill with probability threshold and persistence for each combination of NWP at a fixed lead time range (from 2 to 5.5 days). Each plot represents a different combination of NWP. As a benchmark, the black cross represents the skill of ENS with optimised criteria.

is not suitable as its skill is poorer than the benchmark from day 2 onwards. The current procedure (1D+1P) is also sub-optimal, even though the degradation with lead time is not as severe as in MM.

3) IMPACTS OF PERSISTENCE ON THE WARNING SKILL OF COMBINED NWP

Figure 6 presents the evolution of the gran ensemble skill with probability and persistence for a fixed lead time range (2-5.5 days). The black cross represents the skill of ENS with its optimal warning criteria.

The results reinforce the idea that the persistence criterion limits the maximum attainable skill in ensemble systems. Across all four combinations, the highest skill is reached when persistence is removed (1/1). In this scenario, all combinations exhibit similar maximum skill. Only member

399 weighted (MW) and Brier weighted (BW) marginally surpass the benchmark (ENS). These two
400 combinations show optimal skill within a range of probability thresholds spanning 40% to 70%.

401 The lack of sensitivity to probability is notable in the current criteria, indicated by the dark blue
402 line in the left-hand side pane. Even at the lowest probability threshold, the $f_{0.8}$ score remains at 0.5,
403 unlike other combinations that exhibit a significant decline in performance when the probability
404 threshold takes very low values.

405 4) OPTIMAL WARNING CRITERIA FOR COMBINED NWP

409 Table 4 summarizes the results of the optimization of the warning criteria for the combination
410 methods in an identical way as Table 3 did for the NWP. The results are organised by lead time
411 ranges; for each range, the table presents the optimized criteria and the skill metrics for the four
412 methods. As a benchmark, the first row in each lead time range shows the skill of ENS. A graphical
413 summary of these results can be seen in Appendix C.

414 The benchmark (ENS) is only outperformed in terms of $f_{0.8}$ by some combinations at lead times
415 shorter than 5.5 days. At the shortest lead time range, the inclusion of deterministic NWP is
416 important as they have notable skill (see Figure B1). For this reason, the two combination methods
417 that allocate higher weight to deterministic models (MM and BW) show the highest skill. At lead
418 times from 2 to 5.5 days, BW and MW stand out as the best approaches. They prove comparable
419 skill and optimize similar criteria due to the fact that they give a similar weight to ENS, as it is
420 both the NWP with the largest number of members and with the highest skill. It is worth noting
421 that the combination currently used in EFAS (1D+1P) is the only approach that never outperforms
422 the benchmark. Overall, the optimal grand ensemble can achieve an improvement in skill up to
423 4.2% compared with ENS (up to 10% compared with the current criteria).

424 The optimised criteria show a pattern similar to that observed in the individual NWPs. The
425 persistence criterion is removed in all the combinations at lead times shorter than 7 days. The
426 probability threshold replaces the role of persistence, adopting optimal values notably larger than
427 in the current criteria that decrease with increasing lead time.

428 5) OPTIMAL VERSUS CURRENT WARNING SKILL

435 The maps in Figure 7 illustrate the skill metrics associated with the best grand ensemble —BW
436 with no persistence and a probability threshold of 52.5%— for the lead time range 2 to 5.5 days.

TABLE 4. Summary of the optimization of the warning criteria for the combination methods and four lead time ranges. The initial row serves as a benchmark, indicating the skill of ENS, the top-performing NWP. In each lead time range, bold fonts highlight the highest skill for a specific metric.

lead time	method	probability	persistence	recall	precision	bias	$f_{0.8}$
<2 d	ENS	0.800	1/1	0.681	0.711	0.958	0.699
	1D+1P	0.925	1/1	0.662	0.717	0.923	0.694
	MM	0.775	1/1	0.628	0.838	0.749	0.741
	MW	0.750	1/1	0.673	0.718	0.937	0.700
	BW	0.950	1/1	0.620	0.829	0.748	0.733
2-5.5 d	ENS	0.650	1/1	0.416	0.725	0.574	0.562
	1D+1P	0.600	1/1	0.455	0.643	0.708	0.554
	MM	0.700	1/1	0.424	0.700	0.606	0.559
	MW	0.500	1/1	0.455	0.692	0.658	0.574
	BW	0.525	1/1	0.451	0.700	0.644	0.576
5.5-7 d	ENS	0.450	1/1	0.264	0.605	0.436	0.402
	1D+1P	0.425	1/1	0.262	0.599	0.437	0.399
	MM	0.500	1/1	0.276	0.412	0.670	0.345
	MW	0.425	1/1	0.271	0.579	0.468	0.401
	BW	0.450	1/1	0.268	0.586	0.457	0.400
7-10 d	ENS	0.175	2/2	0.287	0.396	0.725	0.345
	1D+1P	0.250	1/1	0.256	0.414	0.618	0.334
	MM	0.225	2/2	0.252	0.334	0.754	0.296
	MW	0.175	2/2	0.278	0.402	0.692	0.343
	BW	0.300	1/1	0.255	0.445	0.573	0.345

A pronounced contrast is evident in the maps, attributable to the limited number of events in most of the stations, resulting in skill metrics primarily yielding values of either 1 or 0. In terms of $f_{0.8}$, there is a cluster of underperforming points in central Europe (notably the Alps) and central France. The decomposition into recall and precision reveals that the primary contributor to the low f-score is recall, indicating missed events. Conversely, the precision of warnings remains consistently high, with exceptions observed in the Seine and Arno rivers. It is noteworthy the skill in all metrics for the Ebro, Rhine and Meuse—three rivers that experienced significant flood events during the study period.

By applying the optimised criteria, EFAS would be able to predict with at least 2 days lead time 46% of the events, and 71% of all the warnings would be correct. For reference, these figures




FIG. 7. Skill maps for the optimised warning criteria (BW, 52.5% probability and no persistence) at the lead time range from 2 to 5.5 days. The main plot shows the skill in terms of $f_{0.8}$; grey dots are stations for which the f-score cannot be computed due to all instances of hits, misses and false alarms being null. The smaller plots show precision and recall; for the sake of visibility, it only shows the stations for which the specific metric can be computed. In all the plots the direction of the arrows shows the difference in skill between the optimised and the current criteria; positive values imply gains and negative values losses in skill.

represent an improvement in both metrics over the existing criteria, which achieved 41% recall and 62% precision.

6) INFLUENCE OF CATCHMENT AREA ON WARNING SKILL

The preceding sections presented results for a set of stations with a catchment area of at least 2000 km², the current limit in EFAS. To investigate the potential relaxation of this criterion, Figure 8 illustrates the evolution of skill depending on the catchment area limit. The criteria used to create




FIG. 8. Evolution of skill with catchment area limit. Each plot represents a skill metric. In each of the plots, the black line represents the skill of the current warning criteria, while the colour lines depict the skill of the optimized criteria for the 4 combination methods. The vertical, dotted line indicates the catchment area limit currently applied in EFAS.

this figure are those optimized for the fixed area limit of 2000 km² and a lead time range from 2 to 5.5 days (Table 4). The plot compares the skill of the four combination methods (coloured lines) with the skill of the current warning criteria (black line).

The skill of the system improves with catchment area, as expected. However, the skill curves do not increase continuously. Recall experiences a loss in skill for areas ranging from 30,000 to 70,000 km², leading to a decrease in $f_{0.8}$. This loss is caused by missed events in the Guadiana, Seine, Loire, Rhone and Danube catchments. As already mentioned in previous sections, there is a notable gap between recall and precision in both the optimised and current warning criteria, with precision consistently surpassing recall.

The vertical line at 2000 km² illustrates the gain in skill achieved with the optimization, which is sustained across all catchment area values. Reducing the catchment area limit results in minimal loss in $f_{0.8}$, suggesting that this criterion could be relaxed without compromising skill. For instance, the BW combination at 1,000 km² exhibits slightly higher skill ($f_{0.8} = 0.531$) than the current criteria at 2000 km² ($f_{0.8} = 0.518$). The relaxation of the catchment area limit impacts the components of the f-score differently. While precision remains unaffected, there is a loss in recall.

5. Discussion

a. Limitations and moving forward

The definition of the spatial and temporal framework is often a limitation in studies that assess the skill of ensemble flood forecasting systems (Pappenberger et al. 2008b; Cloke and Pappenberger 2009). Due to the rarity of flood events, the statistical robustness of such analysis relies on a limited and often highly correlated sample of events. To address this limitation, long-term and continental or global studies are advocated, such as the one presented here, which spans over a period of two and a half years and extends over a European domain. Nevertheless, our study would benefit from a larger sample of points and flood events. As shown in Figure 7, there are many stations for which the skill metrics cannot be calculated due to null hits, misses or false alarms. A larger sample of observed events would reduce the susceptibility of the analysis to this issue. Two potential solutions emerge: extending the study period, which is limited by the availability of the HEPS forecasts, or using an event threshold of higher recurrence. We decided to work with a sample of stations, instead of all the river cells (Bartholmes et al. 2009), to limit the spatial correlation of events.

In the optimization of the warning criteria, we set a catchment area limit of 2,000 km², a threshold inherited from previous EFAS setups with lower spatial and temporal resolution. The results of this study demonstrate that the current resolution of both NWP and hydrological models can produce comparably skilful warnings at smaller catchments (Figure 8), which can significantly increase the number of stations and flood events analysed (A2). In line with this, the new EFASv5 (JRC 2023) released in September 2023 increased spatial resolution from 5 km to 1 arc-minute, which is expected to further improve the skill of the system in smaller catchments. The skill of this new EFAS version will be analysed as soon as a long enough period or reforecasts are available.

The selection of the target metric is a critical aspect in any skill analysis. Bartholmes et al. (2009) reviewed the desired characteristics of a skill score, and selected odds ratio, HK score and bias as their target metrics. In contrast, we selected the f_β score, a metric commonly used in machine learning, but rarely applied in hydrology. The f_β score has two interesting properties for analysing a flood warning system. Firstly, it is based on the contingency table but avoids using true negatives, which could overestimate skill. Secondly, it allows for tailoring the warning criteria

to the specific use case, as the β parameter can be adjusted to penalize differently misses or false alarms. For instance, Bouttier and Marchal (2024) used the f_2 score to optimize the probability threshold for high-intensity precipitation warnings in France. The β value of 2 penalizes four times more missed events than false alarms, under the assumption that the cost of being hit by an extreme event unprepared is larger than the cost of mobilising resources in a false alarm. In our study, we took the opposite approach and decided to minimize the number of false alarms by using a β value of 0.8. This decision was based on the fact that EFAS complements the national or regional systems by issuing medium-range flood warnings. The aim is to alert authorities, not the general public, only when there is substantial evidence that the event will occur, in order to avoid excessive warnings that could undermine their trust in the system. Additionally, the event threshold (5-year return period) corresponds to a recurrence time shorter than the flood protection levels in many European regions, so we assume that missing minor events will not be as costly as Bouttier and Marchal (2024) argue. In a preliminary analysis, we tested β values of 0.8, 1 and 1.25 and found that the changes in the f-score were minimal (in the order of 3%) and in none of the cases persistence was needed. However, the selection of β affects the optimal probability threshold and the resulting bias, i.e., the distribution of errors between misses and false alarms. Such decision should be taken in collaboration with the end users to tailor the warnings to their specific needs (Lavers et al. 2020; Dasgupta et al. 2023).

b. Addressing study objectives

The first objective of this study was to evaluate the skill of the different NWP within EFAS, with a specific focus on comparing deterministic and probabilistic models. Figure B1 demonstrates that probabilistic NWP outperform deterministic models in terms of Brier skill score across all lead times. Only at short lead times do deterministic NWP reach similar levels of skill. The outcome is similar when applying the warning criteria and measuring warning skill in terms of $f_{0.8}$ (Table 3). At their optimal warning criteria, the probabilistic NWP, particularly ENS, outperform deterministic NWP. The skill of both types of NWP degrades with increasing lead time, casting doubt on their ability to issue flood alerts more than 7 days in advance. The effects of the probability threshold and persistence criteria differs depending on the nature of the NWP. The probability threshold has a substantial effect on the skill of probabilistic models, while it does

not impact deterministic models. Both COS and ENS perform better at probability thresholds over 50%, with a wide range of similarly performing values. On the contrary, persistence proves to be beneficial in improving the warning skill of deterministic NWP, but it limits the skill of probabilistic models. Persistence, a.k.a. poor man's ensemble, is a method of creating an ensemble out of a deterministic forecast (Cloke and Pappenberger 2009). For instance, a persistence of 3/3 is an ensemble of 3 members with a probability threshold of 100%. Our results show that while persistence removes the inherent erratic behaviour of deterministic models and notably improves their skill, it should not be applied to HEPS.

The second objective of this study was to determine the appropriate method to combine several NWP models into a grand ensemble. Careful consideration is necessary in selecting this method to avoid diminishing the skill of the most proficient NWP within the grand ensemble. For instance, the current EFAS warning criteria limit the skill of the system to that of the highest-performing deterministic model (HRES). This limitation is caused by a combination method that excessively relies on deterministic forecasts, as well as the use of persistence. The challenge lies in finding a combination that enhances the skill of the best-performing NWP, which in the case of EFAS is ENS.

Our findings demonstrate that ENS is a very strong baseline, with only a few combination methods yielding gains to its skill. These findings contrast slightly with Pappenberger et al. (2008a), who showed for a specific flood event in Romania in 2007 that grand ensembles can detect extreme events that a single EPS could entirely miss. In our study case, the effectiveness of the grand ensemble is limited to the first 5 days lead time. Overall, the most promising approaches are the member-based (MW) and skill-based (BW) methods. Both methods have similar optimal warning criteria (no persistence and a probability threshold close to 50%), and their success is based on granting most of the weight to ENS. However, the reason why they allocate this weight to ENS differs. In the case of MW, it is coincidental that the most skilful model also possesses the largest number of members. On the contrary, BW grants greater value to ENS because it proved to be a superior model. If a more skilful model with fewer members would be added to the grand ensemble, BW would be able to extract that skill, while MW would not. Therefore, we argue that skill-based combination methods are preferable.

559 *c. Towards operationalization*

560 The decision on the most appropriate combination method cannot be based solely on skill metrics,
561 but also consider other factors such as the operational implementation effort, user comprehensi-
562 bility, or the impact of adding or removing a NWP. Skill-based approaches, such as BW, are
563 theoretically superior as they allocate weights based on quality rather than quantity. They are also
564 the most robust when it comes to adding or removing NWPs. However, they are more challenging
565 to implement and to explain to end users. To derive the matrix of weights (Figure 4), a skill-based
566 approach requires an extensive period of reanalysis and reforecasts upfront, posing a substantial
567 computational burden on system implementation. Additionally, practitioners using the forecasts
568 would require additional training to understand that the ensemble members are not equiprobable,
569 which would be more intuitive.

570 There remains a question on the use of deterministic models in a grand ensemble dominated
571 by a probabilistic NWP. When using a member-based combination, such as MW, deterministic
572 models could be removed from the grand ensemble, as the weight allocated to them is so small
573 that their inclusion barely affects the total probability. Skill-based approaches, instead, allocate
574 larger weights to these models in the very first lead times, when they have proven skill given their
575 often-higher resolution. The inclusion of deterministic models is a topic for future analysis, as the
576 spatial resolution of deterministic and probabilistic NWP are nowadays similar, and centres such
577 as ECMWF plan to discontinue their deterministic models in favour of the ensembles.

578 **6. Conclusions**

579 This study describes a procedure to optimise warning thresholds in HEPS. We investigated the
580 ability of deterministic and probabilistic NWP models in providing correct flood warnings, and we
581 explored the effects of the warning criteria, namely probability threshold and persistence. The final
582 objective of the study is to devise a method to combine several NWP into a grand ensemble that
583 outperforms the most proficient, individual NWP. Our study case is the European Flood Awareness
584 System (EFASv4), a medium-range continental flood warning system.

585 The individual analysis of the NWPs demonstrated that probabilistic models, particularly ENS,
586 provide better flood warnings than deterministic models, even for short lead times, where the
587 latter show some skill. We discovered that the persistence criterion is useful to increase the skill

588 of deterministic models, as it creates a pseudo-ensemble, but is counterproductive in the case
589 of ensembles. The criterion that maximizes ensemble skill is the probability threshold, whose
590 optimum shows a certain degree of uncertainty, but always exceeds a value of 50%.

591 By default, the combination of several NWP into a grand ensemble does not improve the overall
592 skill of the flood warning system. As an example, the current EFAS warning criteria, which uses
593 four NWPs, is outperformed by only using one NWP (ENS) with the adequate criteria. Two of the
594 approaches tested in this work yielded gains in skill compared with the baseline in medium and
595 short lead times. One of these approaches (*member weighted*) assumes equiprobability among all
596 the model runs, and the other one (*Brier weighted*) assigns weights to each NWP based on their
597 probabilistic skill. In our study case, these two approaches yielded similar results because the most
598 proficient NWP is also the one with the largest number of members. Despite that, we argue that
599 the skill-based approach is conceptually more appropriate, gives significance to the inclusion of
600 deterministic NWP in the grand ensemble, and would perform better if the set of NWPs would
601 change. However, other study cases may consider other factors such implementation costs or
602 explainability in the selection of the combination method.

603 Finally, we explored the effects of catchment area in the warning skill, with the idea of relaxing the
604 current minimum catchment area limit. The results indicate that skill deteriorates with decreasing
605 catchment area, as other studies report, but the improvements in meteorological and hydrological
606 models over the last decade allow for a reduction in the minimum catchment area. We have proven
607 that with the optimal criteria we can halve the area limit, therefore cover a larger area, and achieve
608 a similar performance as in the current system setup. We expect that the skill at smaller catchments
609 will be positively affected by the release of EFASv5 with a higher spatial resolution.


610 *Acknowledgments.* The authors would like to thank Jutta Thielen-del Pozo for her revision and
611 contribution to the original manuscript.

612 *Data availability statement.* The EFASv4 reanalysis and forecast discharge data is available at
613 the Climate Data Store of the Copernicus Climate Change Service (<https://cds.climate.copernicus.eu>). All the scripts, pre-processed data and results of the study can be found in this
614 GitHub repository: https://github.com/casadoj/EFAS_skill.
615

616 APPENDIX A

617 Analysis of "observed" events

618 Figure A1 depicts the spatial distribution of the 1239 stations used in the optimization of the
619 warning criteria. The colours represent the amount of "observed" flood events, while the histogram
620 at the bottom shows the frequency of events across the stations.



621 FIG. A1. Geographical distribution of the stations used for optimizing the warning criteria, and number of
622 "observed" flood events during the study period.

623 The optimization of the warning criteria is based on 1239 stations and 874 "observed events".
624 Throughout the study period, 61% of the stations did not exceed the 5-year return period, which

means that hits and misses cannot exist, and the only term in the contingency table will be false alarms. Consequently, the f-score for these stations can be either 0 or null.

The imposition of a minimum contributing area of 2000 km² places a constraint on the number of stations available for optimization, consequently affecting the count of observed events. In the final phase of our analysis, we assessed the implications of this criterion by expanding our consideration to stations with a minimum area of 500 km². Figure A2 represents the relationship between an increasing catchment area threshold and the corresponding number of stations and "observed" events. By excluding the points with contributing area smaller than 2000 km², we discarded 37% of the stations and 48% of the events. The objective of this decision was to use the current warning criteria as benchmark.




FIG. A2. Number of stations (orange) and observed events (blue) by catchment area threshold. The vertical, dotted line represents a catchment area of 2000 km², the current limit in the EFAS warnings.

APPENDIX B

Probabilistic skill of individual NWP

In the initial assessment of NWP skill, we calculated the Brier score for each NWP at daily lead times. The Brier score, as an error metric, yields values that can be challenging to interpret, given the dependence of error magnitude on the specific variable —particularly in the context of rare events like floods, where the magnitude is inherently low. Instead, we employ the Brier skill score

(BSS) to gauge the relative skill of a model in comparison to a reference. Models outperforming the reference receive a positive BSS, while those underperforming register a negative BSS. For the reference, we opted for a dummy model that never predicts an event ($P_{pred} = 0$) (Legg and Mylne 2004).




FIG. B1. Brier skill score of the four NWP used in EFAS at daily lead times. The reference ($BSS = 0$) is a model that never predicts an event. Blue lines represent probabilistic models, and orange lines deterministic ones.

The plot above demonstrates that probabilistic models exhibit greater skill than deterministic ones. Only within very short lead times (0-1 days) do deterministic models approximate the skill of probabilistic models. This is particularly significant for EFAS warnings, as warnings are issued only at lead times longer than 2 days (indicated by the leftmost dotted line). As lead time increases, there is a degradation in skill. This decline impacts ENS to a lesser extent than the other models. Both deterministic models display poor skill at their forecast horizon.

APPENDIX C

Roebber diagram: a visual summary of the results

The Roebber diagram in Figure C1 offers an alternative perspective on the results of the optimization of the warning criteria. Roebber diagrams condense four skill metrics into a single plot, with an ideal model positioned at the top right corner (Roebber 2009). The figure C1 shows simultaneously the skill of the current warning criteria, the optimization of individual NWP (Table 3), and the optimization of combination methods (Table 4) for the lead time range 2-5.5 days.




FIG. C1. Roebber diagram of the optimized criteria at the lead time range from 2 to 5.5 days. Four skill metrics are represented: precision and recall in the X and Y axis, bias by the dashed lines, and the $f_{0.8}$ score by the solid lines. Triangles illustrate individual NWP and circles the combination of NWP. For reference, the skill of the current warning criteria is shown as a black cross.

All NWP models and the current criteria exhibit negative bias, indicating that EFAS underpredicts the occurrence of flood events. In the case of the optimizations carried out in this analysis, the selection of the $f_{0.8}$ score as the target metric in the optimization forces this bias, as it prioritizes precision over recall. It is remarkable that the improved $f_{0.8}$ of the probabilistic models is a result of higher precision, while recall values remain similar across all NWPs and the current criteria.

References

Alfieri, L., B. Bisselink, F. Dottori, G. Naumann, A. D. Roo, P. Salamon, K. Wyser, and L. Feyen, 2017: Global projections of river flood risk in a warmer world. *Earth's Future*, **5** (2), 171–182, <https://doi.org/10.1002/2016EF000485>.

Alfieri, L., F. Pappenberger, F. Wetterhall, T. Haiden, D. Richardson, and P. Salamon, 2014: Evaluation of ensemble streamflow predictions in europe. *Journal of Hydrology*, **517**, 913–922, <https://doi.org/10.1016/j.jhydrol.2014.06.035>.

Bartholmes, J., J. Thielen, M. Ramos, and S. Gentilini, 2009: The european flood alert system efas – part 2: Statistical skill assessment of probabilistic and deterministic operational forecasts. *Hydrology and Earth System Sciences*, **13**, 141–153, <https://doi.org/10.5194/hess-13-141-2009>.

- Bartholmes, J., and E. Todini, 2005: Coupling meteorological and hydrological models for flood forecasting. *Hydrology and Earth System Sciences*, **9** (4), 333–346, <https://doi.org/10.5194/hess-9-333-2005>.
- Bauer, P., A. Thorpe, and G. Brunet, 2015: The quiet revolution of numerical weather prediction. *Nature*, **525** (7567), 47–55, <https://doi.org/10.1038/nature14956>.
- Bouttier, F., and H. Marchal, 2024: Probabilistic short-range forecasts of high precipitation events: optimal decision thresholds and predictability limits. *EGUsphere [preprint]*, <https://doi.org/10.5194/egusphere-2023-3111>.
- Brier, G., 1950: Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, **78**.
- Buizza, R., 2008: The value of probabilistic prediction. *Atmospheric Science Letters*, **9** (2), 36–42, <https://doi.org/10.1002/asl.170>.
- Burek, P., J. van der Knijff, and A. D. Roo, 2013: Lisflood. distributed water balance and flood simulation model. Joint Research Centre - European Commission, <https://doi.org/10.2788/24719>.
- Caretta, M., and Coauthors, 2023: Water. *Climate Change 2022 – Impacts, Adaptation and Vulnerability*, H. Pörtner, D. Roberts, M. Tignor, E. Poloczanska, K. Mintenbeck, A. Alegría, M. Craig, S. Langsdorf, S. Löschke, V. Möller, A. Okem, and B. Rama, Eds., Cambridge University Press, 551–712, <https://doi.org/10.1017/9781009325844.006>.
- Cloke, H., and F. Pappenberger, 2009: Ensemble flood forecasting: A review. *Journal of Hydrology*, **375**, 613–626, <https://doi.org/10.1016/j.jhydrol.2009.06.005>.
- Cloke, H., F. Pappenberger, P. Smith, and F. Wetterhall, 2017: How do i know if i’ve improved my continental scale flood early warning system? *Environmental Research Letters*, **12**, <https://doi.org/10.1088/1748-9326/aa625a>.
- Dasgupta, A., and Coauthors, 2023: Connecting hydrological modelling and forecasting from global to local scales: Perspectives from an international joint virtual workshop. *Journal of Flood Risk Management*, <https://doi.org/10.1111/jfr3.12880>.

De Roo, A., C. Wesseling, and W. V. Deursen, 2000: Physically based river basin modelling within a gis: The lisflood model. *Hydrological Processes*, **14**, 1981–1992, [https://doi.org/10.1002/1099-1085\(20000815/30\)14:11/12<1981::aid-hyp49>3.0.co;2-f](https://doi.org/10.1002/1099-1085(20000815/30)14:11/12<1981::aid-hyp49>3.0.co;2-f).

De Roo, A., and Coauthors, 2003: Development of a european flood forecasting system. *International Journal of River Basin Management*, **1 (1)**, 49–59, <https://doi.org/10.1080/15715124.2003.9635192>.

Demeritt, D., S. Nobert, H. Cloke, and F. Pappenberger, 2013: The european flood alert system and the communication, perception, and use of ensemble predictions for operational flood risk management. *Hydrological Processes*, **27**, 147–157, <https://doi.org/10.1002/hyp.9419>.

Diomede, T., and Coauthors, 2008: Discharge prediction based on multi-model precipitation forecasts. *Meteor. Atmos. Phys.*, **101 (3-4)**, 245–265, <https://doi.org/10.1007/s00703-007-0285-0>.

EEA, 2023: Economic losses from weather- and climate-related extremes in Europe. Tech. rep., European Environmental Agency - European Union. URL <https://www.eea.europa.eu/en/analysis/indicators/economic-losses-from-climate-related?activeAccordion=ecdb3bcf-bbe9-4978-b5cf-0b136399d9f8>.

Gupta, H., H. Kling, K. Yilmaz, and G. Martínez, 2009: Decomposition of the mean squared error and nse performance criteria: Implications for improving hydrological modelling. *Journal of Hydrology*, **377**, 80–91, <https://doi.org/10.1016/J.JHYDROL.2009.08.003>.

Hanssen, A., and W. Kuipers, 1965: *On the Relationship Between the Frequency of Rain and Various Meteorological Parameters*. Koninklijk Nederlands Meterologisch Instituut. Mededelingen en Verhandelingen, Staatsdrukkerij- en Uitgeverijbedrijf.

JRC, 2020: European Flood Awareness System v4.0. Joint Research Centre - European Commission, URL <https://confluence.ecmwf.int/display/CEMS/EFAS+v4.0>.

JRC, 2023: European Flood Awareness System v5.0. Joint Research Centre - European Commission, URL <https://confluence.ecmwf.int/display/CEMS/Latest+operational+EFAS+release>.

Kling, H., M. Fuchs, and M. Paulin, 2012: Runoff conditions in the upper Danube basin under an ensemble of climate change scenarios. *Journal of Hydrology*, **424-425**, 264–277, <https://doi.org/10.1016/j.jhydrol.2012.01.011>.

- Knoben, W., J. Freer, and R. Woods, 2019: Technical note: Inherent benchmark or not? comparing nash-sutcliffe and kling-gupta efficiency scores. *Hydrology and Earth System Sciences*, **23**, 4323–4331, <https://doi.org/10.5194/hess-23-4323-2019>.
- Lavers, D., and Coauthors, 2020: A vision for hydrological prediction. *Atmosphere*, **11** (3), <https://doi.org/10.3390/atmos11030237>.
- LeClerc, J., and S. Joslyn, 2015: The cry wolf effect and weather-related decision making. *Risk Analysis*, **35** (3), 385–395, <https://doi.org/10.1111/risa.12336>.
- Legg, T., and K. Mylne, 2004: Early Warnings of Severe Weather from Ensemble Forecast Information. *Wea. Forecasting*, **19** (5), 891–906.
- Molteni, F., R. Buizza, T. Palmer, and T. Petroliagis, 1996: The ECMWF Ensemble Prediction System: Methodology and validation. *Quart. J. Roy. Meteor. Soc.*, **122**, 73–119, <https://doi.org/10.1002/qj.49712252905>.
- Nevo, S., and Coauthors, 2022: Flood forecasting with machine learning models in an operational framework. *Hydrology and Earth System Sciences*, **26** (15), 4013–4032, <https://doi.org/10.5194/hess-26-4013-2022>.
- Pagano, T., and Coauthors, 2014: Challenges of operational river forecasting. *J. Hydrometeor.*, **15**, 1692–1707, <https://doi.org/10.1175/jhm-d-13-0188.1>.
- Pappenberger, F., J. Bartholmes, J. Thielen, H. Cloke, R. Buizza, and A. D. Roo, 2008a: New dimensions in early flood warning across the globe using grand-ensemble weather predictions. *Geophys. Res. Lett.*, **35** (10), <https://doi.org/10.1029/2008GL033837>.
- Pappenberger, F., K. Beven, N. Hunter, P. Bates, B. Gouweleeuw, J. Thielen, and A. D. Roo, 2005: Cascading model uncertainty from medium range weather forecasts (10 days) through a rainfall-runoff model to flood inundation predictions within the European Flood Forecasting System (EFFS). *Hydrology and Earth System Sciences*, **9** (4), 381–393, <https://doi.org/10.5194/hess-9-381-2005>.
- Pappenberger, F., M. Ramos, H. Cloke, F. Wetterhall, L. Alfieri, K. Bogner, A. Mueller, and P. Salamon, 2015: How do I know if my forecasts are better? Using benchmarks in hydrological

ensemble prediction. *Journal of Hydrology*, **522**, 697–713, <https://doi.org/10.1016/j.jhydrol.2015.01.024>.

Pappenberger, F., K. Scipal, and R. Buizza, 2008b: Hydrological aspects of meteorological verification. *Atmospheric Science Letters*, **9**, 43–52, <https://doi.org/10.1002/asl.171>.

Paprotny, D., A. Sebastian, O. Morales-Nápoles, and S. Jonkman, 2018: Trends in flood losses in Europe over the past 150 years. *Nature Communications*, **9** (1), <https://doi.org/10.1038/s41467-018-04253-1>.

Ramos, M., J. Bartholmes, and J. T. del Pozo, 2007: Development of decision support products based on ensemble forecasts in the european flood alert system. *Atmospheric Science Letters*, **8**, 113–119, <https://doi.org/10.1002/asl.161>.

Roebber, P., 2009: Visualizing multiple measures of forecast quality. *Wea. Forecasting*, **24**, 601–608, <https://doi.org/10.1175/2008WAF2222159.1>.

Roulin, E., and S. Vannitsem, 2005: Skill of Medium-Range Hydrological Ensemble Predictions. *J. Hydrometeor.*, **6** (5), 729–744, <https://doi.org/10.1175/JHM436.1>.

Thielen, J., J. Bartholmes, M. Ramos, and A. D. Roo, 2009a: The european flood alert system – part 1: Concept and development. *Hydrology and Earth System Sciences*, **13**, 125–140, <https://doi.org/10.5194/hess-13-125-2009>.

Thielen, J., K. Bogner, F. Pappenberger, M. Kalas, M. del Medico, and A. D. Roo, 2009b: Monthly-, medium-, and short-range flood warning: Testing the limits of predictability. *Meteor. Appl.*, **16**, 77–90, <https://doi.org/10.1002/met.140>.

Wetterhall, F., and Coauthors, 2013: HESS Opinions ”Forecaster priorities for improving probabilistic flood forecasts”. *Hydrology and Earth System Sciences*, **17**, 4389–4399, <https://doi.org/10.5194/hess-17-4389-2013>.

WMO, 2021: WMO Atlas of mortality and economic losses from weather, climate and water extremes (1970-2019). Tech. rep., World Meteorological Organization, Geneva.

788 Yang, S., T. Yang, Y. Chang, C. Chen, M. Lin, J. Ho, and K. Lee, 2020: Development of
789 a hydrological ensemble prediction system to assist with decision-making for floods during
790 typhoons. *Sustainability*, **12** (10), <https://doi.org/10.3390/su12104258>.